

Voice-over · 12 clips

Record each paragraph as its own clip and name the files **01** to **12** (Voice Memos or QuickTime: File > New Audio Recording; any format is fine). If you stumble, record that clip again; nothing else needs redoing.

Speak calmly and leave half a second of silence at the start and end; I trim it and match the video to the length of each clip, so there is no timing to follow.

Quiet room, phone or laptop mic 20 cm away, same distance for every clip.

Clip 01

On screen: Slide 1 · cover

Hi, I'm Diego, and this is my entry to the Factored AI and Data Hackathon. It is a dispute agent for LATAM Bank that knows when not to act. It handles card disputes end to end, in Spanish and Portuguese, and hands off to a person when policy says so.

Clip 02

On screen: Slide 2 · why disputes

We chose disputes from the data. Complaint contacts are fixed on first contact only 44 percent of the time, against 92 percent for transactional ones. Unrecognized charges and undue fees are about 40 percent of formal complaints. And every transaction ties cleanly to its product and customer, so the agent can verify the exact charge.

Clip 03

On screen: Live app, scenario list

Here it is, live, on the deployed app. The interface is in Spanish, because that is what LATAM Bank customers speak. On the right, demo scenarios with real sandbox customers. On top, a progress bar that follows the agent's real stage.

Clip 04

On screen: Ambiguous scenario in Portuguese: document and SMS code

A customer writes in Portuguese about a charge they don't recognize. The router detects Portuguese and a dispute. Before looking at any data, the agent verifies identity: first the document, then a one-time code, simulated here as an SMS.

Clip 05

On screen: The agent lists the matching charges; the customer answers

There are several charges at that merchant, so instead of guessing, it asks. The customer answers in their own words, and the model maps that to the right option.

Clip 06

On screen: A question at the confirmation

Before any write, it asks for an explicit yes. A question is not a yes: it asks again, and writes nothing.

Clip 07

On screen: Dispute filed, card blocked, live trace

Now it files the dispute, verifies it by reading it back, and offers to block the card. The progress bar names the outcome. And the live trace shows every tool call, the policy rule, and the cost of each turn: fractions of a cent.

Clip 08

On screen: Large amount: hand-off to a person

And when it should not act: a large withdrawal. After verifying identity, policy says a person must review it, so the agent hands off with a reference. The advisor receives everything already verified, and the progress bar marks the confirmation step as skipped.

Clip 09

On screen: Slide 3 · architecture

The core design decision: the model understands, code decides. A learned router and an LLM read the message. The LLM never sees the document or the one-time code. A state machine, a versioned policy, and customer-scoped bank tools decide and act. When the router and the model disagree, the agent asks the customer instead of guessing.

Clip 10

On screen: Slide 4 · learned router

The router is the learned component. It was chosen by a rule fixed before the sealed, hand-written test set was opened. Both learned routers beat the keyword baseline by more than 20 points of macro F1.

Clip 11

On screen: Slide 5 · held-out results

End to end, we ran 230 conversations through both systems, on a fresh set of records sealed before the run. The hybrid safely resolved every in-scope dispute, 104 out of 104, against 86.5 percent for rules. It made every required handoff, with zero unsafe outcomes. The version before had one, and the slide shows how we got here.

Clip 12

On screen: Slide 6 · route to production

It runs today, with tracing, bounded retries, a spending cap, and a safe fallback. And we list what is missing to make it real. Build something that works, prove it works, and know when not to act. Thank you.